Local LLM Performance on Windows 11

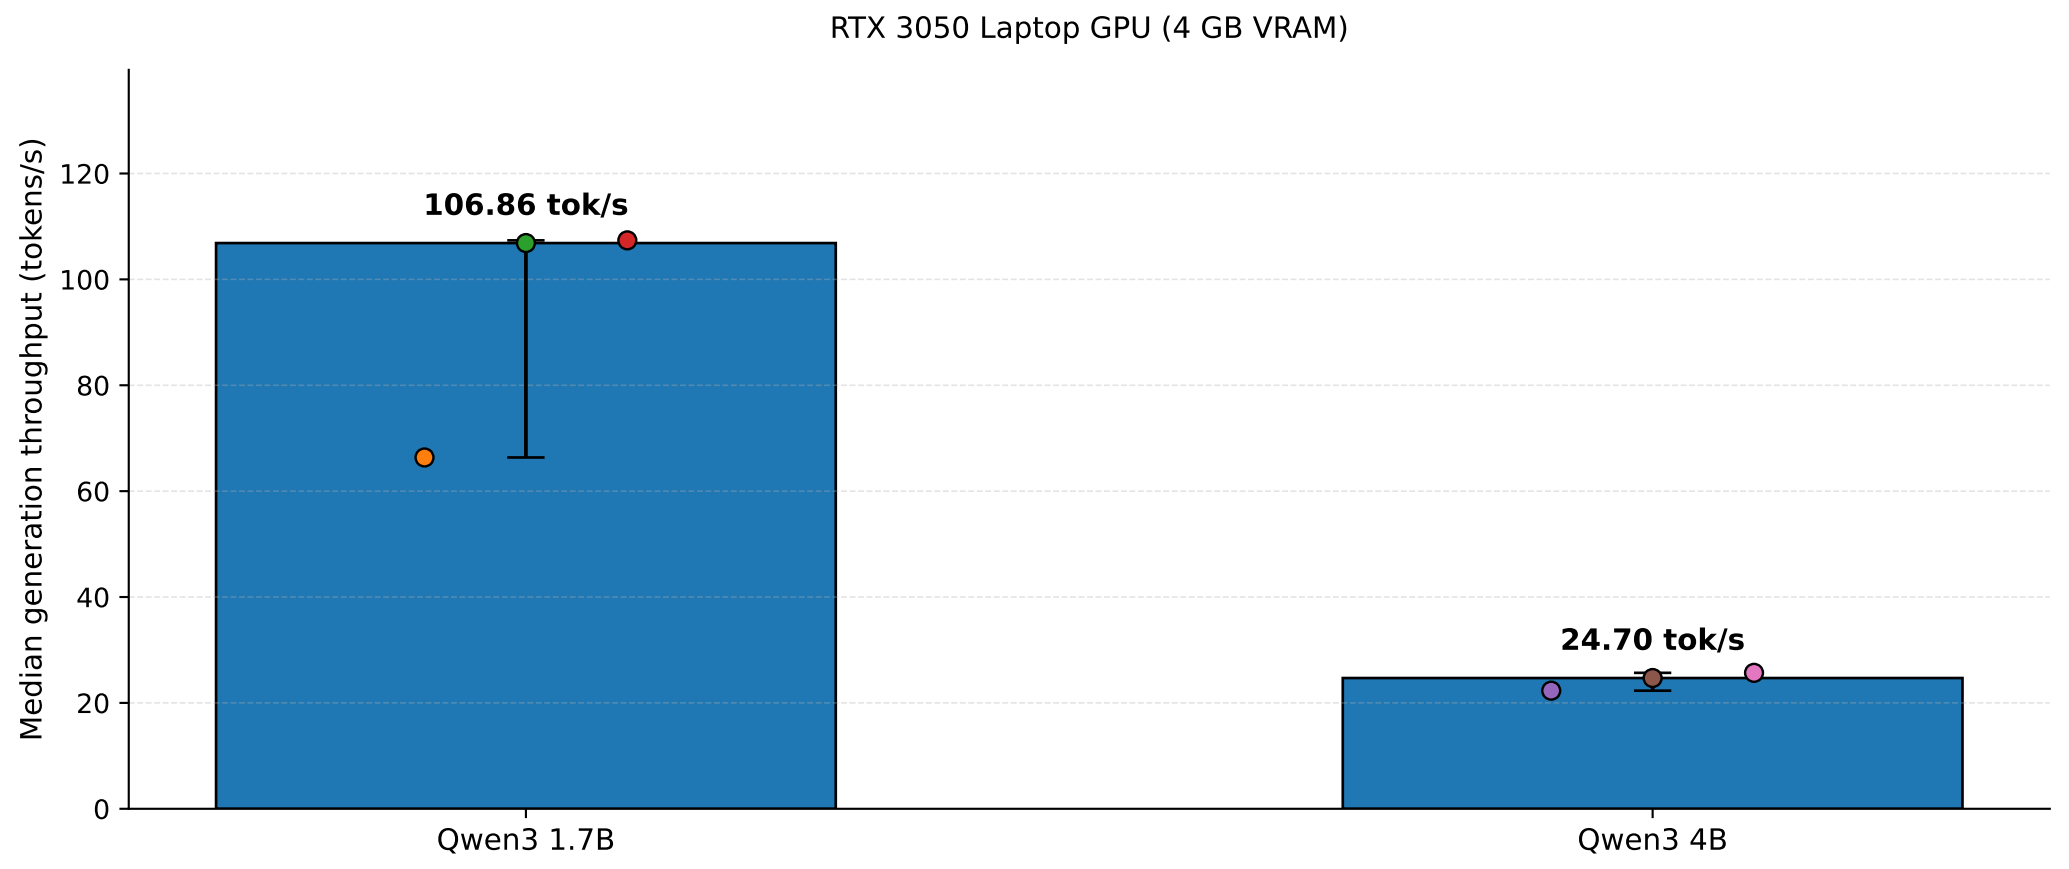

Qwen3 1.7B
RAM Δ 0.67 GB
VRAM Δ 1716 MiB
Execution 100% GPU
Range 66.36–107.37 tok/s

Qwen3 4B
RAM Δ 1.50 GB
VRAM Δ 2312 MiB
Execution 33% CPU / 67% GPU
Range 22.31–25.68 tok/s

Ollama 0.34.4 · context 4096 · 128 requested tokens · temperature 0 · n = 3 runs/model
Observed median throughput ratio: Qwen3 1.7B ≈ 4.33× Qwen3 4B