Security Boundary and Agent-Tool Behavior

Controlled evaluation of workspace boundaries and agent tool invocation

A. Filesystem boundary behavior

Tool	Controlled operation	Observed result	Interpretation
ReadFileTool	Relative + absolute outside-workspace read	**BLOCKED**	Workspace boundary enforced
WriteFileTool	Relative + absolute outside-workspace write	**BLOCKED**	Workspace boundary enforced
BashTool	Outside-workspace read + write	**ALLOWED**	Normal host permissions
PythonTool	Outside read: tested Outside write: not tested	**READ ALLOWED**	Write behavior not claimed

Boundary interpretation: dedicated ReadFile/WriteFile tools enforced the tested workspace boundary. Bash and Python should not be treated as operating-system filesystem sandboxes.

B. Agent tool-invocation reliability

Scenario	Result	Observation
Qwen3 1.7B natural read task	**FAIL**	No actual tool call was emitted
Qwen3 4B natural read task	**FAIL**	No actual tool call was emitted
Forced valid tool protocol	**PASS**	One valid fenced tool block was parsed
read_file backend execution	**PASS**	Tool completed with exit code 0
Forced end-to-end read	**PASS**	Returned synthetic marker ORBIT-GLASS-219

Interpretation: natural task failures indicate protocol-emission unreliability in this controlled test; they do not imply that Qwen3 models generally cannot use tools.

Frontend observation: the shell UI appeared disabled while an Agent-intent request contained allow_bash=true. End-to-end shell execution from that disabled UI state was not demonstrated.